Project Scope

Team 3 - Governance for Ethical AI Application

Background
Organization is a large public university based in California, which has observed rapid escalation of AI use across faculty, staff, and students. The Uni has implemented a basic set of guidelines about data use in prompts and syllabus language. However, it is considering a much more systematic approach, driven by various stakeholders.

Scope
The scope is to develop a governance framework for safe and ethical AI. Consultant will present insights and propose recommendations that the U leadership team can use as input for formulating the ethical AI strategy. Key required activities include:

- Establish a framework and methodology for the project.
- Research the industry to understand what U's competitors are doing.
- Research adjacent industries to see what can be borrowed from other market leaders.
- Develop recommendations:
  - Develop an assessment of commercially available approaches, software, etc., that provide potential building blocks.
  - Propose new work practices, training, learning and development, to accompany a potential new ethical AI governance framework.
  - Discuss risks and costs, as well as benefits.
- Outline next steps for the U to pursue its quest to develop compelling new ethical AI framework.
- Provide an executive level report and presentation to the U's management team.

SoW specifications:
- From the perspective of the university
- Goal: finding a middle ground to both use AI and still foster learning
- Think: what's a good way to standardize? Where is the line?
- "Adjacent industries" research — maybe look into large companies like Google and retail clothing companies
  - phrasely.ai?
- Know the largest competitors in AI (ChatGPT, Claude, Copilot, etc.)
- We're SOW #3

Ideas:
- Make a web framework?
  - Important categories: investors, students, faculty, future employers
  - List assumptions and definitions (think BA200)
  - Consider different subjects/topics

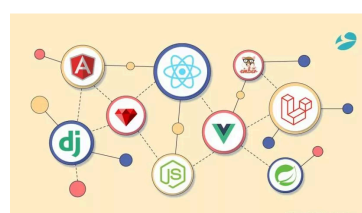

- Website Deliverable
  - Sample syllabus
  - Research
  - Discussion forum

- Have at least one case from a company and one case from a school to include in the final report?
- Cost estimates and timelines
    - Training, policies, etc.
-

Library Consultation:
- Look how AI works, what its doing– helps us understand elements of ethical AI in particular to the university
    - Ex: AI picking up on biases and prejudices without us knowing
- Umich has set up their own standards like the Generative AI Guidance
- Gives us a couple closed data chat gpt vibed tools
- AIethicist.org is rly helpful
- Gives us a summary of the guidance of the top ten universities

Research

Gantt Chart:
https://docs.google.com/spreadsheets/d/16IuSwuE6qRhA_cPlbpY5yHyCRuPNtbNdiMZyHyjvfGQ/edit?gid=1115838130#gid=1115838130

Kresge Library Services Page: https://kresgeguides.bus.umich.edu/TO-300-2026-Winter-B

Colleges: Due Sunday

- Considerations
  - Political standing
  - Size: pick larger
  - Region
  - public/private
- Maybe find class syllabus examples if possible? @ different universities
- Look at library sources
- Ideas:
  - MIT,
  - UC Berkley,
  - UT Austin,
  - Arizona State

Jiya: UCLA, NYU, Stanford

Noah:

Ian:

Nacha: UCB, MIT, UT Austin

Connor:

Nadia: UChicago, Purdue, UIUC (?), UMICH

Nadia's Research

Syllabus Examples: Syllabus Examples

University of Chicago

- https://genai.uchicago.edu/en/resources/students
    - In favor of AI adoption within reason
    - Can be used for learning but is limited
    - Consider privacy, ethical considerations, etc.
    - Classes and events for learning what AI tools UChicago offers
- https://genai.uchicago.edu/about/generative-ai-guidance
- Specific list of allowances: https://genai.uchicago.edu/generative-ai-tools/approved-and-restricted-ai-tools
- Awaiting syllabi

Purdue University

- https://www.purdue.edu/vpec/policies/information-technology/viia5/
    - AI can increase productivity but must be monitored
    - Protect sensitive and restricted data
    - University-wide policy but leaves discretion to unit heads as well
    - Ensure accuracy
    - Can be used to inform but not make decisions
- https://it.purdue.edu/ai/ai-review-governance/?_gl=1*1khlmr2*_gcl_au*MTEwMzI5MDA2Mi4xNzc0MjIxNDgx*_ga*OTk3ODM1ODA1LjE3NzQyMjE0ODE.*_ga_PF1CYQ27F6*czE3NzQyODE5OTYkbzIkZzEkdDE3NzQyODIwMzckajE5JGwwJGgw
    - Review process for AI tools
- Received syllabus

University of Illinois at Urbana-Champaign

- Reviewed two syllabi from a student; neither mentioned AI specifically, but both mentioned citing your work or attribution of sources
- https://ldlprogram.web.illinois.edu/academic-integrity-statement/
    - Provide prompt, output, and a note of your experience (advantages, disadvantages, etc.)
- Seemed pretty non-specific
- https://www.vpaa.uillinois.edu/digital_risk_management/generative_ai

University of Michigan Ann Arbor

- https://genai.umich.edu/resources/faculty/course-policies
    - Discretion by course; either encouraged, allowed, or forbidden
    - Requires transparency when allowed
- Will add syllabi to the folder

Nacha's Research

University of California Berkley

- Students and staff access to AI tools such as Google Gemini
    - https://technology.berkeley.edu/news/new-ai-tools-students
- UCB launched a "Campus AI Sandbox" for secure testing, and promotes AI-related research and education
    - https://ai.berkeley.edu/resources/campus-ai-sandbox
- BAIR(Berkeley AI Research): flagship lab that is famous for “Generalist” AI
    - Berkeley Lab uses AI to accelerate experiments in materials science, chemistry, and climate science, including designing artificial enzymes
    - https://bair.berkeley.edu/blog/
    - https://www.lbl.gov/research/ai/
- Students and staff can use Zoom AI Companion for meeting summaries and in-meeting questions
    - https://technology.berkeley.edu/news/new-ai-tools-students
- Have a bunch of AI research which is human-centered like helping restore voices, predicting weather, etc
    - https://www.berkeley.edu/ai/#:~:text=Real%2Dworld%20Impact,medical%20breakthroughs%20to%20climate%20change

Massachusetts Institute of Technology

- Schwarzman College of Computing:
    - Founded with a $1 billion commitment
    - "Hub" for all things AI: designed to break down silos, pairing computer scientists with philosophers and urban planners
    - https://computing.mit.edu/
- CSAIL: largest research lab at MIT. It focuses on everything from "Liquid Neural Networks" (AI that can change its structure on the fly) to generative biology
    - https://www.csail.mit.edu/
- AI & Creativity Hub: deepen ties between computing and design as advances in artificial intelligence are reshaping how ideas are conceived and shared
    - https://news.mit.edu/2026/mit-hasso-plattner-institute-collaborative-hub-for-ai-and-creativity-0320
- MIT leads global conversations on how governments should regulate AI, focusing on actionable "Policy Briefs" rather than just vague ethical theories
    - https://aipolicyforum.mit.edu/
- They have integrated "Social and Ethical Responsibilities of Computing" (SERC) into the core curriculum for every engineering student
    - https://computing.mit.edu/cross-cutting/social-and-ethical-responsibilities-of-computing/serc-scholars-program/

University of Texas Austin

- UT’s Grand Challenge: A 10-year research initiative focused on "Ethical AI" in areas like surveillance, smart cities, and disinformation
    - https://bridgingbarriers.utexas.edu/good-systems
- Machine Learning Laboratory: Funded by major alumni and works on the foundations of AI like the math and algorithms that make models faster and cheaper to run
    - https://ml.utexas.edu/
- UT Sage: AI tutor trained on specific course materials
    - Used in classrooms to provide personalized support
    - https://tech.utexas.edu/initiatives/artificial-intelligence
- UT Spark offers a secure, campus-wide AI platform for chatbots and content generation
    - https://tech.utexas.edu/initiatives/artificial-intelligence
- The UT REAL Health AI initiative, in partnership with Qualified Health, is implementing HIPAA-compliant AI tools for clinical use cases
    - https://www.utsystem.edu/sites/health-intelligence-platform/blog/ut-real-health-ai-launch-2025-11-19#:~:text=The%20initiative%20has%20also%20partnered%20with%20Qualified,Workgroup**%20Develops%20training%20programs%20and%20educational%20resources
- UT leads the NSF-funded Institute for Foundations of Machine Learning, focusing on creating AI that can learn from less data
    - https://ai.utexas.edu/news/2025/ut-lead-next-gen-ai

Have gathered syllabus from:
- College of charleston
- Texas A&M
- NC State
- UVA

Ian

ASU

First U.S. university to partner with OpenAI, every one of their 181,000 students, faculty, and staff gets free ChatGPT Edu (GPT-5) access

Negotiated a privacy agreement where student data is never used to train OpenAI's models

Faculty set their own AI policies per course and must state them in the syllabus

The university actively encourages disclosure and citation over prohibition

They officially discourage AI detection tools, believing the tools are too unreliable to be fair to students

The library has guides on how to properly cite AI

Built a "Principled Innovation" framework that was reviewed and signed off by legal, cybersecurity, and the general counsel before rollout, not just an edu department deciding on their own

https://ai.asu.edu/

Connor's Research

Stanford:
-Prevents Graduate level classes from banning AI on any take home assignment.
    -Highly encourages students to use AI, and has creative prompts attached to assignments to get students to cite AI and reflect upon its usefulness
-They have their own AI tools for sensitive class material
-Access to Gemini
https://tlhub.stanford.edu/docs/course-policies-on-generative-ai-use/
-Their AI policy was deleted from their community standards website
-Most decisions around AI seem to be deferred to the instructors, although not having AI policy listed is sus

University of Virginia
-Gives professors the discretion to create their own standards
-Standards much more relaxed–They don't have a universal honor code with guidelines and punishments
-Professors can restrict AI use in any way in class, however, from an administrative side, they are encouraged to create clear guidelines and allow AI in appropriate use cases
-Their policies about AI are designed more like FAQ rather than a large, standardized document
https://genai.provost.virginia.edu/guidance-for-faculty-students
https://honor.virginia.edu/resources/artificial-intelligence

Jiya's Research

UCLA:
https://dts.ucla.edu/initiatives/ai/ai-use-recommendation-guide
https://dts.ucla.edu/initiatives/ai/guiding-principles-responsible-use
https://teaching.ucla.edu/special-initiatives/generative-ai-in-teaching-learning/

Key Points

- Responsible AI Use: UCLA emphasizes that AI should be a tool to augment, not replace, human thought, particularly in academic work.
- Data Security: Users must protect sensitive data. The official recommendation is to use enterprise-supported AI platforms rather than consumer versions to ensure data protection.
- Ethical Considerations: Systems should be evaluated for bias, fairness, and potential inaccuracies before implementation, with a focus on transparency and ethical, human-centric application.
- Academic Integrity: Students must cite AI writing tools (e.g., ChatGPT) if used. Using AI to generate content without proper attribution or turning in AI-generated text as one's own is considered academic dishonesty.
- Governance and Oversight: The UCLA Health AI Council (HAIC) is developing risk-management frameworks to evaluate technologies from procurement through operation.

Stanford
https://uit.stanford.edu/security/responsibleai
https://ucomm.stanford.edu/policies-and-guidance/ai-guidelines-marketing-and-communications
https://news.stanford.edu/stories/2025/01/report-outlines-stanford-principles-for-use-of-ai

- Emphasizes data privacy, don't input sensitive or confidential information into AI tools
- Encourages awareness, caution, and transparency when using AI
- AI should support, not replace, human work
- Users are fully responsible for accuracy and outcomes of AI-generated content
- Some uses (e.g., official chatbots) require institutional approval
- Promotes experimentation with guardrails, not strict bans
- Focuses on balancing innovation with ethical responsibility
- Highlights risks like plagiarism, misuse, and unclear authorship
- Uses a flexible, principle-based approach rather than strict universal rules

NYU
https://www.nyu.edu/faculty/teaching-and-learning-resources/Student-Learning-with-Generative-AI.html
https://guides.nyu.edu/ai-tools
https://steinhardt.nyu.edu/faculty-and-staff/academic-affairs/steinhardt-ai-hub/academic-integrity-and-syllabus-support-age

- Academic Integrity: Use of generative AI must align with instructor policies. Submitting AI-generated work as one's own constitutes a violation.

- Verification: All AI-generated content must be independently verified for accuracy, bias, and fabricated citations.
- Privacy & Data Security: Never input confidential, proprietary, or student-identifiable information into public AI tools.
- AI Detection: NYU cautions against relying on AI detection tools, acknowledging their unreliability.
- Sustainability: NYU encourages sustainable AI use, recommending efficient usage to mitigate the energy consumption of large language models (LLMs).
- Generative AI Services: NYU provides access to authorized services, such as Copilot, accessible via university login.
- NYU Langone Health: Focuses on AI in medicine through principles of patient safety, quality, and efficiency.
- Research & Governance: The NYU Stern Center for Business and Human Rights studies risks like bias, misinformation, and hallucinations.
- AI Policy Hub: A forum focused on AI governance in public safety.

Noah's Research

University of North Carolina at Chapel Hill (UNC)

- Encourages transparency and disclosure of AI use in academic work
- Focus on risks such as bias, hallucinations, and misinformation in AI outputs
- UNC School of Data Science and Society integrates AI, ethics, and public policy into curriculum and research
  https://datascience.unc.edu/
- Carolina AI Initiative promotes interdisciplinary AI research across healthcare, social sciences, and public policy
  https://ai.unc.edu/
- UNC emphasizes AI literacy and responsible adoption rather than restriction

University of California Los Angeles (UCLA)

- UCLA encourages instructors to define their own AI policies and integrate AI into coursework responsibly
- UCLA DataX initiative focuses on data science + AI collaboration across disciplines
  https://datax.ucla.edu/

University of California Santa Cruz (Arts-Focused UC)

- UCSC Center for Research in Open Source Software + AI collaboration
  https://cross.ucsc.edu/
- Arts + Media integration of AI through Digital Arts and New Media (DANM) program
  https://danm.ucsc.edu/
- AI research in computational media, games, and creative technologies
  https://engineering.ucsc.edu/departments/computational-media/

Key Takeaways

- No university outright outlaws AI usage
- Majority leave it up to professor and departmental discretion
- Biggest concerns is not with efficiency of AI but rather reliability
  - AI has major concerns regarding bias answers and hallucinations

Team Charter

Team 5 Charter

Norms and expectations for TO 300 Section 008 Team 3

<u>Group Members</u>

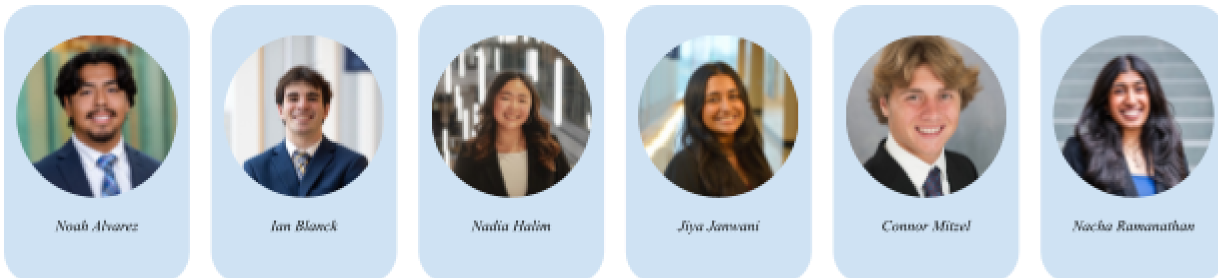

<u>Group Goals</u>

- Maintain open, respectful, and timely communication (respond to messages within 24 hours) so all members are informed and aligned.
- Complete assigned tasks on time and at high quality to avoid last-minute work, holding one another accountable through regular check-ins.
- Start and end meetings on time with a clear agenda and action items.
- Represent the team professionally in presentations and written work.

<u>Meetings</u>

Before

- Before meetings, we plan to set a weekly recurring meeting time that best fits all of our schedules. This will make coordination a lot easier throughout the semester and ensure that all members of our team are consistently aware of the meeting time.
- Meetings will generally take place online through Google Meet unless otherwise specified.

During

- During the meetings, we will focus on high-level tasks to use our time efficiently and also do tasks that require team deliberations. Getting this done in person supports our team's effort to work cohesively opposed to completing work that does not mesh well together. We will also delegate specific tasks during this time and make a game plan for what each task consists of and when a reasonable time to have the task done is.
- We will make a list of tasks to be completed before the next meeting as well as who is responsible for each task to avoid confusion or miscommunication.

After

- After our meeting, we will check to ensure that nothing has come up during the time of the next meeting and text our group chat with any clarifying questions about the meeting

that we have. It is very important that everyone has a good understanding of their role and what is expected of them before the next meeting. We will work on any needed work for the project before the next meeting to ensure we stay on track to complete the project.
- Tasks will be assigned based on either the preferences or the strengths of each group member. Each member will be responsible for taking on a task that matches the workload of other team members and completing it by an agreed deadline.

General Expectations
- Each member comes to every meeting with an understanding of the meeting's agenda and is ready to share any work that was to be completed before the meeting. We will listen respectfully to each other, engage in project discussions, and offer constructive feedback during meetings.

Absences
- Absences from a meeting due to other events or commitments must be communicated at least 24 hours in advance. Emergencies should be communicated immediately.
- Any work assigned to the absent team member should be communicated to them by the end of the meeting, and the absent teammate must have it completed by the agreed deadline. Any clarifying questions or concerns should be communicated through the group chat.

Communication Methods

Preferred Methods
- Group members will answer texts in the iMessage group chat in a timely manner. To ensure communication flows and questions are answered, we may follow up on unanswered messages with additional messages.

General Use
- The group message will be used to schedule meetings, assign work, give updates on assignments, etc.

Emergencies
- If the situation calls for some level of urgency and a team member isn't answering group or direct messages, a phone call may be used.

Communication Style

Norms & Expectations
- Group members will respect each other and their individual ideas.
- Everyone should speak in a respectful manner to other team members to ensure a smooth work environment during meetings.
- If someone is not pulling their weight, the team should communicate their concerns immediately to the member rather than immediately contacting outside sources. If this method does not work, then additional group members or course leadership may be

involved. However, we will make every effort to first communicate directly with group members about concerns.

Accountability

- Team members are expected to contribute their fair share to the project. If conflicts prevent a group member from completing their work, transparent communication to the group is expected to ensure the project will still be completed.
- Everyone should be accountable for taking on an equal workload and completing it by the agreed deadline. This includes managing time accordingly to ensure the work is completed, rather than employing excuses that may result in the group falling behind.

Handling Conflict

- When the team disagrees, we will use the disagreement method of a mediator. The members in disagreement should call upon another group member to hear both sides and determine which one is the best route to pursue so the group can best succeed.
- If the conflict is group-wide, an outside neutral mediator (such as the professor or the SSA) will be briefed on the conflict and asked for advice on how to proceed.
- In general, group members should be flexible in their thinking and willing to compromise for the good of the group, allowing ideas and the end deliverable to be at their strongest.

Executive Update

Executive Update Team 3
3/19/25

Understanding of the SOW
The central issue, outlined in the SOW, is the rapid adoption of AI without regard for ethical standards in education. Our main objective is to provide a framework for students and professors to follow when using AI in the classroom. To create this framework, we must first research the ethical standards of similar institutions and adjacent industries. When determining our recommendation, we should consider the current software and frameworks in place, the adoption process, and the associated risks and benefits. Our final report should outline the University's next steps in ethical AI adoption.

Meeting Schedule
Tentatively, we have scheduled a recurring meeting time on Mondays at 12 pm-1 pm. The meetings will be online to allow for easy access unless otherwise noted in the meeting invitation. If an in-person meeting is necessary, we will communicate the location in the meeting invitation and our text group chat.

Kresge Library Discussion
Meeting with our librarian helped us learn more about which sources would be the best for us to leverage when researching to create the foundation of our framework. Through the library discussion with Laura, we were able to gather content that helped us expand our knowledge and framework such as the case studies of guidance at the Big Ten schools which enabled us to want to explore further and try to gather real-time information and syllabi from those colleges and other ones as well.

Roles and Responsibilities
We have created a team charter to guide our interactions and divide roles between our group members. We understand that we are all responsible for the outcome of this project, and as such, we all must be involved in the project's successful completion.

Issues
Our only issue that needs the professor's help at this early point in the project would be integrating class material. Some of what we have learned isn't as applicable to our project, and we are interested about how to integrate it going forward.

Gantt Chart and Project Plan
For our final project, we want to create three key deliverables: a sample website, a sample syllabus, and the final presentation. Our desired timeline is linked here.

Deliverables

Gantt Chart:
https://docs.google.com/spreadsheets/d/16IuSwuE6qRhA_cPlbpY5yHyCRuPNtbNdiMZyHyjvfGQ/edit?gid=1115838130#gid=1115838130

Canva Site: https://canva.link/efpwon3ckdys16d

Website Outline

Home Page
- Define AI
- Purpose of site
- Latest news
- FAQ

Students
- Using AI to learn
  - Fact-checking
- Resources
  - Just a “this is where to list resources” or “example 1” chart or something like that
  - Purpose is to provide students with helpful tools?

Faculty
- Training modules summary?
- How to craft an AI syllabus
  - “Prohibited” example
  - “Allowed” example
  - “Encouraged” example
- Resources

Tools
- Database of approved AI tools from the university

About
- Purpose of site
- Encouraging AI as a tool and opportunity to enhance productivity but not at the cost of learning

Contact
- Submit requests for info or questions
  - Some sort of department that should be handling this?
- AI committee?

Action Items

Week of (3/30)

Website: Nadia, Jiya

- Rough draft done by Friday

Syllabus: Ian, Nacha

- Collect all the syllabi + maybe add into a drive
- Rough draft done by Friday

Report: Connor, Noah

- Outline done by Friday
- Do we need more research about certain areas?
- What's cost looking like? Can we estimate?

Everyone

- Review deliverables and add comments by Sunday

Week of (3/23)

Syllabus: Ian and Nacha

- Figure out the layout
- Get syllabi from other schools

Website: Nadia and Jiya

- Begin structuring site layout and functionality
- Look at other universities' GenAI sites for guidance
- Research different tab ideas

Report: Noah

-